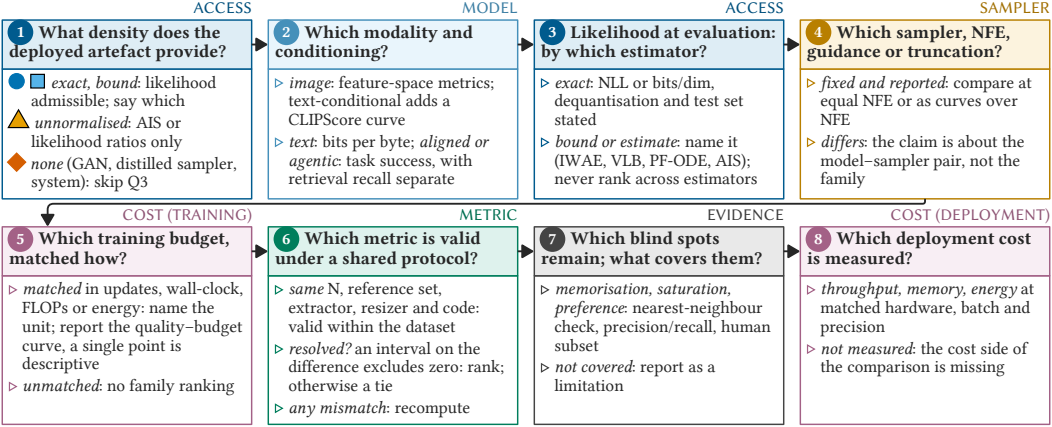
ACCESS
1 What density does the deployed artefact provide?
exact, bound: likelihood admissible; say which
unnormalised: AIS or likelihood ratios only
none (GAN, distilled sampler, system): skip Q3
MODEL
2 Which modality and conditioning?
image: feature-space metrics; text-conditional adds a CLIPScore curve
text: bits per byte; aligned or agentic: task success, with retrieval recall separate
ACCESS
3 Likelihood at evaluation: by which estimator?
exact: NLL or bits/dim, dequantisation and test set stated
bound or estimate: name it (IWAE, VLB, PF-ODE, AIS); never rank across estimators
SAMPLER
4 Which sampler, NFE, guidance or truncation?
fixed and reported: compare at equal NFE or as curves over NFE
differs: the claim is about the model–sampler pair, not the family
COST (TRAINING)
5 Which training budget, matched how?
matched in updates, wall-clock, FLOPs or energy: name the unit; report the quality–budget curve, a single point is descriptive
unmatched: no family ranking
METRIC
6 Which metric is valid under a shared protocol?
same N, reference set, extractor, resizer and code: valid within the dataset
resolved? an interval on the difference excludes zero: rank; otherwise a tie
any mismatch: recompute
EVIDENCE
7 Which blind spots remain; what covers them?
memorisation, saturation, preference: nearest-neighbour check, precision/recall, human subset
not covered: report as a limitation
COST (DEPLOYMENT)
8 Which deployment cost is measured?
throughput, memory, energy at matched hardware, batch and precision
not measured: the cost side of the comparison is missing